


Convolutions

IMAGE PROCESSING WITH KERAS IN PYTHON




Ariel Rokem

Senior Data Scientist, University of
Washington

Using correlations in images

- Natural images contain spatial correlations
- For example, pixels along a contour or edge
- How can we use these correlations?

Biological inspiration



What is a convolution?

```
array = np.array([0, 0, 0, 0, 0, 1, 1, 1, 1, 1])
kernel = np.array([-1, 1])
conv = np.array([0, 0, 0, 0, 0, 0, 0, 0, 0])
conv[0] = (kernel * array[0:2]).sum()
conv[1] = (kernel * array[1:3]).sum()
conv[2] = (kernel * array[2:4]).sum()
...
for ii in range(8):
    conv[ii] = (kernel * array[ii:ii+2]).sum()
conv
```

```
array([0, 0, 0, 0, 1, 0, 0, 0, 0])
```

Convolution in one dimension

```
array = np.array([0, 0, 1, 1, 0, 0, 1, 1, 0, 0])
kernel = np.array([-1, 1])

conv = np.array([0, 0, 0, 0, 0, 0, 0, 0, 0])
for ii in range(8):
    conv[ii] = (kernel * array[ii:ii+2]).sum()

conv
```

```
array([ 0,  1,  0, -1,  0,  1,  0, -1,  0])
```

Image convolution

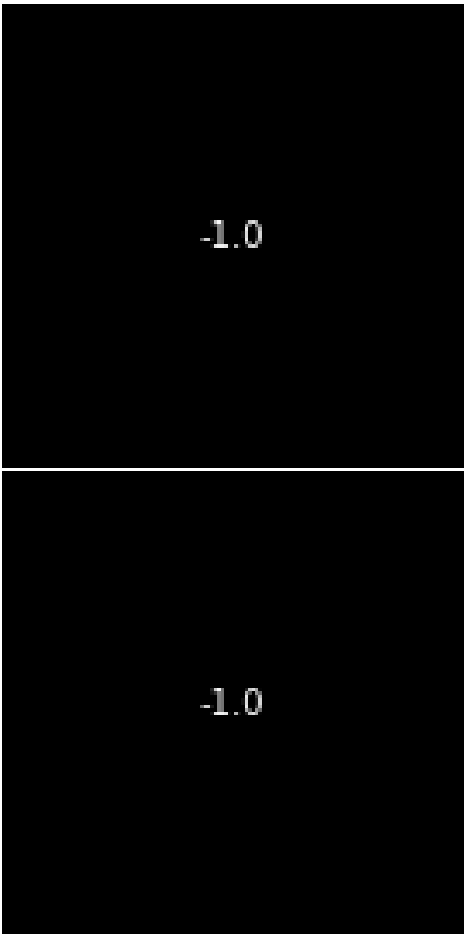





Image convolution



Two-dimensional convolution

```
kernel = np.array([[ -1,  1],  
                   [ -1,  1]])  
  
conv = np.zeros((27, 27))  
  
for ii in range(27):  
    for jj in range(27):  
        window = image[ii:ii+2, jj:jj+2]  
        conv[ii, jj] = np.sum(window * kernel)
```


Convolution



Let's practice!


IMAGE PROCESSING WITH KERAS IN PYTHON

Implementing convolutions in Keras

IMAGE PROCESSING WITH KERAS IN PYTHON

Ariel Rokem

Senior Data Scientist, University of
Washington



Keras Convolution layer

```
from keras.layers import Conv2D  
Conv2D(10, kernel_size=3, activation='relu')
```

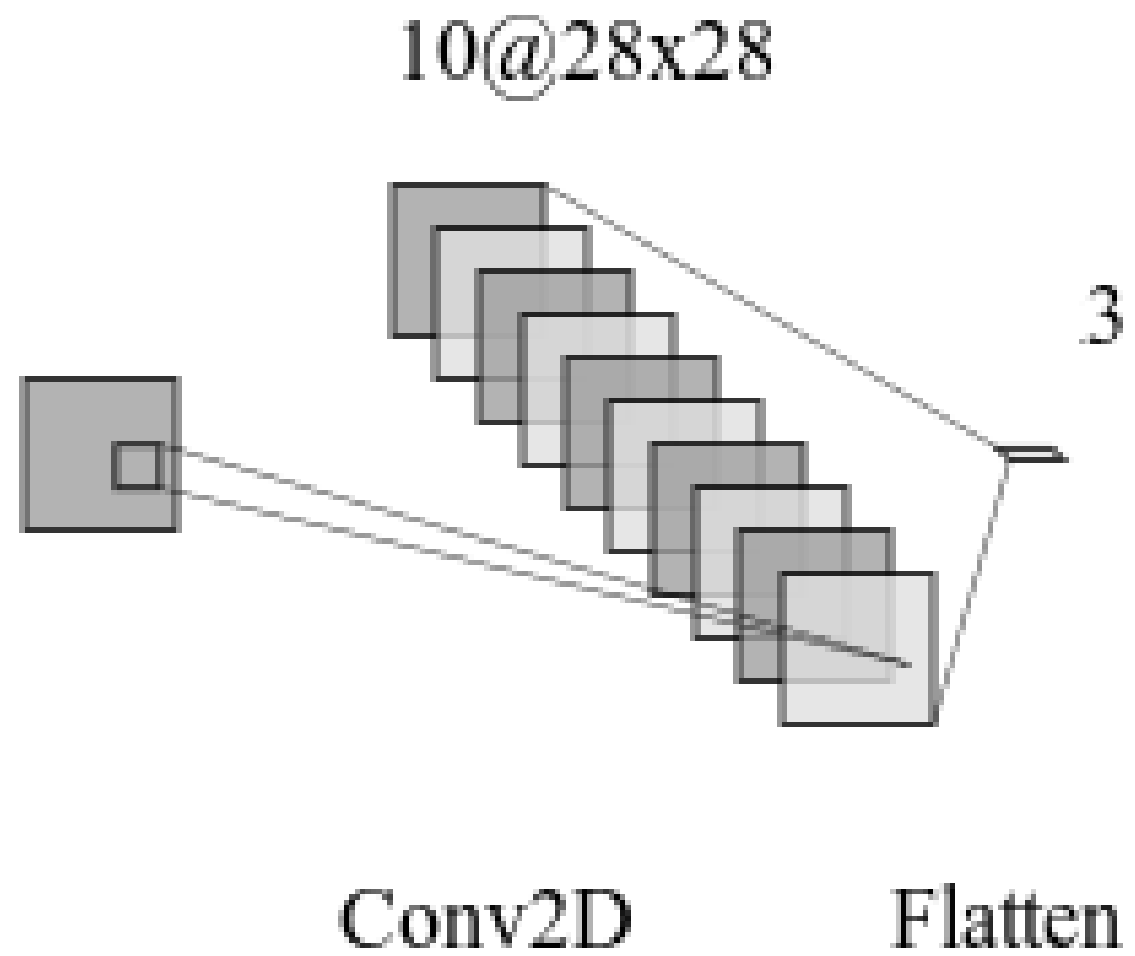
Integrating convolution layers into a network

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten

model = Sequential()
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)))

model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Our CNN



Fitting a CNN

```
model.compile(optimizer='adam',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

```
train_data.shape
```

```
(50, 28, 28, 1)
```

```
model.fit(train_data, train_labels, validation_split=0.2,  
          epochs=3)
```


```
model.evaluate(test_data, test_labels, epochs=3)
```

Let's practice!

IMAGE PROCESSING WITH KERAS IN PYTHON

Tweaking your convolutions


IMAGE PROCESSING WITH KERAS IN PYTHON




Ariel Rokem

Senior Data Scientist, University of Washington

Convolution



Convolution with zero padding




Zero padding in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',  
                 input_shape=(img_rows, img_cols, 1)),  
            padding='valid')
```


Zero padding in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',  
                 input_shape=(img_rows, img_cols, 1)),  
           padding='same')
```

Strides




Strides in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',  
                 input_shape=(img_rows, img_cols, 1)),  
            strides=1)
```

Strides in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',  
                 input_shape=(img_rows, img_cols, 1)),  
            strides=2)
```

Example



Calculating the size of the output

$$O = ((I - K + 2P)/S) + 1$$

where


- I = size of the input
- K = size of the kernel
- P = size of the zero padding
- S = strides

Calculating the size of the output

$$28 = ((28 - 3 + 2)/1) + 1$$

$$10 = ((28 - 3 + 2)/3) + 1$$

Dilated convolutions




Dilation in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',  
                 input_shape=(img_rows, img_cols, 1)),  
            dilation_rate=2)
```

Going deeper


IMAGE PROCESSING WITH KERAS IN PYTHON



Ariel Rokem

Senior Data Scientist, University of
Washington


Network with one convolutional layer



Network with one convolutional layer: implementation

```
model = Sequential()  
model.add(Conv2D(10, kernel_size=2, activation='relu',  
                 input_shape=(img_rows, img_cols, 1)))  
model.add(Flatten())  
model.add(Dense(3, activation='softmax'))
```


Building a deeper network



Building a deep network

```
model = Sequential()
model.add(Conv2D(10, kernel_size=2, activation='relu',
                 input_shape=(img_rows, img_cols, 1),
                 padding='equal'))

# Second convolutional layer
model.add(Conv2D(10, kernel_size=2, activation='relu'))


model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Why do we want deep networks?





Convolution
Pooling
Softmax
Other



Features in early layers



Features in intermediate layers



Features in late layers



How deep?


- Depth comes at a computational cost
- May require more data

Let's practice!

IMAGE PROCESSING WITH KERAS IN PYTHON

How many parameters?

IMAGE PROCESSING WITH KERAS IN PYTHON



Ariel Rokem

Senior Data Scientist, University of
Washington

Counting parameters

```
model = Sequential()

model.add(Dense(10, activation='relu',
                input_shape=(784,)))

model.add(Dense(10, activation='relu'))

model.add(Dense(3, activation='softmax'))
```

```
# Call the summary method
model.summary()
```

```
-----
Layer (type)                 Output Shape              Param #
=====
dense_1 (Dense)              (None, 10)                7850
-----
dense_2 (Dense)              (None, 10)                110
-----
dense_3 (Dense)              (None, 3)                 33
=====
Total params: 7,993
Trainable params: 7,993
Non-trainable params: 0
-----
```

Counting parameters

```
model.add(Dense(  
    10, activation='relu',  
    input_shape=(784,)))
```

$$\begin{aligned} \text{parameters} &= 784 * 10 + 10 \\ &= 7850 \end{aligned}$$

```
model.add(Dense(  
    10, activation='relu'))
```

$$\begin{aligned} \text{parameters} &= 10 * 10 + 10 \\ &= 110 \end{aligned}$$

```
model.add(Dense(  
    3, activation='softmax'))
```

$$\begin{aligned} \text{parameters} &= 10 * 3 + 3 \\ &= 33 \end{aligned}$$

$$7850 + 110 + 33 = 7993$$

```
model.summary()
```

```
-----  
Layer (type)                 Output Shape                 Param #  
-----  
dense_1 (Dense)              (None, 10)                   7850  
-----  
dense_2 (Dense)              (None, 10)                   110  
-----  
dense_3 (Dense)              (None, 3)                    33  
-----  
Total params: 7,993  
Trainable params: 7,993  
Non-trainable params: 0  
-----
```

The number of parameters in a CNN

```
model = Sequential()

model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(28, 28, 1), padding='same'))

model.add(Conv2D(10, kernel_size=3, activation='relu',
                 padding='same'))

model.add(Flatten())

model.add(Dense(3, activation='softmax'))
```

```
model.summary()
```

```
-----  
Layer (type)                 Output Shape              Param #  
-----  
conv2d_1 (Conv2D)            (None, 28, 28, 10)       100  
-----  
conv2d_2 (Conv2D)            (None, 28, 28, 10)       910  
-----  
flatten_3 (Flatten)          (None, 7840)              0  
-----  
dense_4 (Dense)              (None, 3)                 23523  
-----  
Total params: 24,533  
Trainable params: 24,533  
Non-trainable params: 0  
-----
```

The number of parameters in a CNN

```
model.add(
    Conv2D(10, kernel_size=3,
           activation='relu',
           input_shape=(28, 28, 1),
           padding='same'))

model.add(
    Conv2D(10, kernel_size=3,
           activation='relu',
           padding='same'))

model.add(Flatten())
```

```
model.add(Dense(
    2, activation='softmax'))
```

$$\begin{aligned} \text{parameters} &= 9 * 10 + 10 \\ &= 100 \end{aligned}$$

.

$$\begin{aligned} \text{parameters} &= 10 * 9 * 10 + 10 \\ &= 910 \end{aligned}$$

$$\text{parameters} = 0$$

$$\begin{aligned} \text{parameters} &= 7840 * 3 + 3 \\ &= 23523 \end{aligned}$$

Increasing the number of units in each layer

```
model = Sequential()

model.add(Dense(5, activation='relu',
                input_shape=(784, ), padding='same'))

model.add(Dense(15, activation='relu', padding='same'))

model.add(Dense(3, activation='softmax'))
```

```
model.summary()
```

```
-----  
Layer (type)                 Output Shape              Param #  
-----  
dense_1 (Dense)              (None, 5)                 3925  
-----  
dense_2 (Dense)              (None, 15)                90  
-----  
dense_3 (Dense)              (None, 3)                 48  
-----  
Total params: 4,063  
Trainable params: 4,063  
Non-trainable params: 0  
-----
```

Increasing the number of units in each layer

```
model = Sequential()

model.add(Conv2D(5, kernel_size=3, activation='relu',
                 input_shape=(28, 28, 1),
                 padding="same"))

model.add(Conv2D(15, kernel_size=3, activation='relu',
                 padding="same"))

model.add(Flatten())

model.add(Dense(3, activation='softmax'))
```

```
model.summary()
```


```
-----  
Layer (type)                 Output Shape                 Param #  
-----  
conv2d_12 (Conv2D)           (None, 28, 28, 5)           50  
-----  
conv2d_13 (Conv2D)           (None, 28, 28, 15)          690  
-----  
flatten_6 (Flatten)          (None, 11760)               0  
-----  
dense_9 (Dense)              (None, 3)                   35283  
-----  
Total params: 36,023  
Trainable params: 36,023  
Non-trainable params: 0  
-----
```

Let's practice!

IMAGE PROCESSING WITH KERAS IN PYTHON

Reducing parameters with pooling

IMAGE PROCESSING WITH KERAS IN PYTHON



Ariel Rokem


Senior Data Scientist, University of
Washington


```
model.summary()
```


```
-----  
Layer (type)                 Output Shape                 Param #  
-----  
conv2d_12 (Conv2D)           (None, 28, 28, 5)           50  
-----  
conv2d_13 (Conv2D)           (None, 28, 28, 15)          690  
-----  
flatten_6 (Flatten)          (None, 11760)               0  
-----  
dense_9 (Dense)              (None, 3)                   35283  
-----  
Total params: 36,023  
Trainable params: 36,023  
Non-trainable params: 0  
-----
```
















Implementing max pooling

```
result = np.zeros((im.shape[0]//2, im.shape[1]//2))  
result[0, 0] = np.max(im[0:2, 0:2])  
result[0, 1] = np.max(im[0:2, 2:4])  
result[0, 2] = np.max(im[0:2, 4:6])
```

...

```
result[1, 0] = np.max(im[2:4, 0:2])  
result[1, 1] = np.max(im[2:4, 2:4])
```

...

Implementing max pooling

```
for ii in range(result.shape[0]):  
    for jj in range(result.shape[1]):  
        result[ii, jj] = np.max(im[ii*2:ii*2+2, jj*2:jj*2+2])
```


Max pooling in Keras

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten, MaxPool2D

model = Sequential()
model.add(Conv2D(5, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)))

model.add(MaxPool2D(2))

model.add(Conv2D(15, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)))

model.add(MaxPool2D(2))

model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

```
model.summary()
```

```
-----
Layer (type)                 Output Shape              Param #
=====
conv2d_1 (Conv2D)            (None, 26, 26, 5)        50
-----
max_pooling2d_1 (MaxPooling2 (None, 13, 13, 5)        0
-----
conv2d_2 (Conv2D)            (None, 11, 11, 15)       690
-----
max_pooling2d_2 (MaxPooling2 (None, 5, 5, 15)        0
-----
flatten_1 (Flatten)          (None, 375)              0
-----
dense_1 (Dense)              (None, 3)                1128
=====
```